

Post-mortem memory leak diagnosis with ELF and Virtual Functions

Peter Edwards, Arista
peadar@arista.com

Scene:

- Large, long-running process, serving HTTP requests.
- C-with-classes
- No memory diagnostics designed in to speak of
- Customer running a few thousand transactions per second through process
- Process slowly leaking memory
- Cumulative effect meant process would vomit up a core after a day or so.
- Not reproducible in test

Where's the hole in the bucket?

- Could potentially provide an instrumented build, but:
 - conservative customers won't always agree to this easily
 - may reduce performance below acceptable threshold
- Core file is the best thing we have to work with
- There's always 'strings'
 - `$ cat core | strings | sort | uniq -c | sort -n`

C++ vtables

For classes with virtual functions:

- There's one virtual function table
- Each instance of that type starts with a pointer to the table
- The compiler will generate a symbol for the table, eg:
 - “vtable for Namespace::Type”
- Mangled for the linker:
 - `_ZTVN9Namespace4TypeE`

C++ vtables - concept

```
class Base {  
public:  
    virtual void f();  
    virtual void g() = 0;  
    virtual void h();  
    virtual ~Base();  
};
```

```
class Derived : public Base {  
public:  
    virtual void f();  
    virtual void g();  
    virtual ~Derived();  
};
```

C++ vtables - generated vtable

`_ZTV7Derived:`

```
.quad 0  
.quad _ZTI7Derived  
.quad _ZN7Derived1fEv  
.quad _ZN7Derived1gEv  
.quad _ZN4Base1hEv  
.quad _ZN7DerivedD1Ev  
.quad _ZN7DerivedD0Ev
```

C++ vtables - generated table, thru c++filt

vtable for Derived:

```
.quad 0  
.quad typeid for Derived  
.quad Derived::f()  
.quad Derived::g()  
.quad Base::h()  
.quad Derived::~~Derived()  
.quad Derived::~~Derived()
```

vptrs and vtables



Setting vptrs

- Calling a virtual function when constructing or destructing a base type calls the base type's virtual functions, as for destruction
- Implemented by overwriting the vptr to the more derived type after the base is constructed, and vice versa on destruct
- Each instance of a live object has a pointer to its vtable
- These 32- or 64- bit values are unlikely to occur otherwise.

Core files

```
$objdump -x ./corefile
```

```
....
```

Program Header:

```
NOTE off 0x0000000000000740 vaddr 0x0000000000000000 paddr 0x0000000000000000 align 2**0
      filesz 0x0000000000000e64 memsz 0x0000000000000000 flags ---
LOAD off 0x0000000000002000 vaddr 0x0000000000400000 paddr 0x0000000000000000 align 2**12
      filesz 0x0000000000001000 memsz 0x0000000000003000 flags r-x
LOAD off 0x0000000000003000 vaddr 0x0000000000602000 paddr 0x0000000000000000 align 2**12
      filesz 0x0000000000001000 memsz 0x0000000000001000 flags r--
LOAD off 0x0000000000004000 vaddr 0x0000000000603000 paddr 0x0000000000000000 align 2**12
      filesz 0x0000000000001000 memsz 0x0000000000001000 flags rw-
LOAD off 0x0000000000005000 vaddr 0x0000000000246500 paddr 0x0000000000000000 align 2**12
      filesz 0x0000000000003200 memsz 0x0000000000003200 flags rw-
LOAD off 0x00000000000037000 vaddr 0x00007fa04b603000 paddr 0x0000000000000000 align 2**12
      filesz 0x0000000000001000 memsz 0x000000000001c0000 flags r-x
```

Finding our memory leak

- Any leaked object on the heap that has virtual functions will have a telltale pointer at the start of the allocated memory to the vtbl for that type
- Process:
 - Find all the symbols in the image that look like vtables, i.e.,
 - Walk the entire core, 4 bytes at a time, looking for pointers to the vtables

Demo

Questions

<https://github.com/peadar/pstack>

<https://github.com/peadar/pstack/blob/master/canal.cc>